CONFIDENTIAL

Last updated: Aug 19, 2025, 8:01 PM

# Factuality Side-by-Side Evaluation Instructions

Version: Feb 14, 2025

Overall, given the multiple questions and required research, a given task may require quite some time. **Take the time necessary to do a thorough search**. This can be up to in general 30-45 minutes, or even over an hour for particularly difficult cases. **REMEMBER - You are fact checking every single fact until and unless you find a fully Inaccurate result!**

Note - this set of instructions is for versions of this task both **with** and **without** search snippets.


- Factuality Side-by-Side Evaluation Instructions
- AI Response Factuality Side-by-Side Overview
  - What is Factuality?
- Ratings
  - Overall Questions
  - Response Questions
  - Overall Comparison (Only for Understandable prompts)
  - Examples of identifying Factual Claims
  - How to Research
  - Reputable Sources
- FAQ
  - What should my comments look like?
  - Are these models up to date?
  - Do I still rate the Responses even if they are a bad match to the Prompt?
  - Do I still rate the Responses even if the Prompt is nonsense or not in English?
  - (New Example Added) What do I do with time-dependent searches or time-specific answers?
  - Do I need to fact-check links?
  - What do I do with searches for rapidly changing time-dependent searches (weather, prices, etc?)
  - How do I interpret the “Search results” when they are present?
  - If the model is asked to summarize, am I fact-checking the information in the summary?
  - Are links claims?
  - When should I mark a task as “Task Cannot be Rated” (rather than Skip)?
  - Can you expand what you mean by “cannot assess the prompt/response pair”?
  - What should I do when the tables in the response are overlapping?
  - I don’t have a “Which response is better?” button, what should I do?
  - If all claims are Accurate, what should I write in the “Briefly explain your rating for Response A/B” box?
  - Can I see some more examples of “hedging”?
  - What do hints look like? What do I do if the hint is wrong?
- Examples
  - Examples of Accuracy ratings
  - Examples of hedging
  - Examples of Good and Bad comments


# AI Response Factuality Side-by-Side Overview

In this task, you will be provided with a **User Prompt** and two **AI Responses**. You may see some **Time** and **Location** context about when the **Response** was generated. In addition, you may be given a **Hint** about what the correct response is or should look like. You may have access to **search results** the model had available to help create its responses (beyond its normal training data)

You will first determine if the **Task is ratable,** and if so, answer **two initial questions** about each Response. If a Response

contains Factual information, you will need to **perform research to assess whether that information is correct**. You will note if any incorrect information is specifically bad on certain axes. At the end, if the **Prompt is understandable,** you will **directly compare** the two Responses and **explain why you made your decision.**

**Note!:** You are only judging the information in the responses. You are not judging any information in the search results (other than as needed to validate information in the Responses). Notably, we don't care if any of the links in the search results are any good or not.

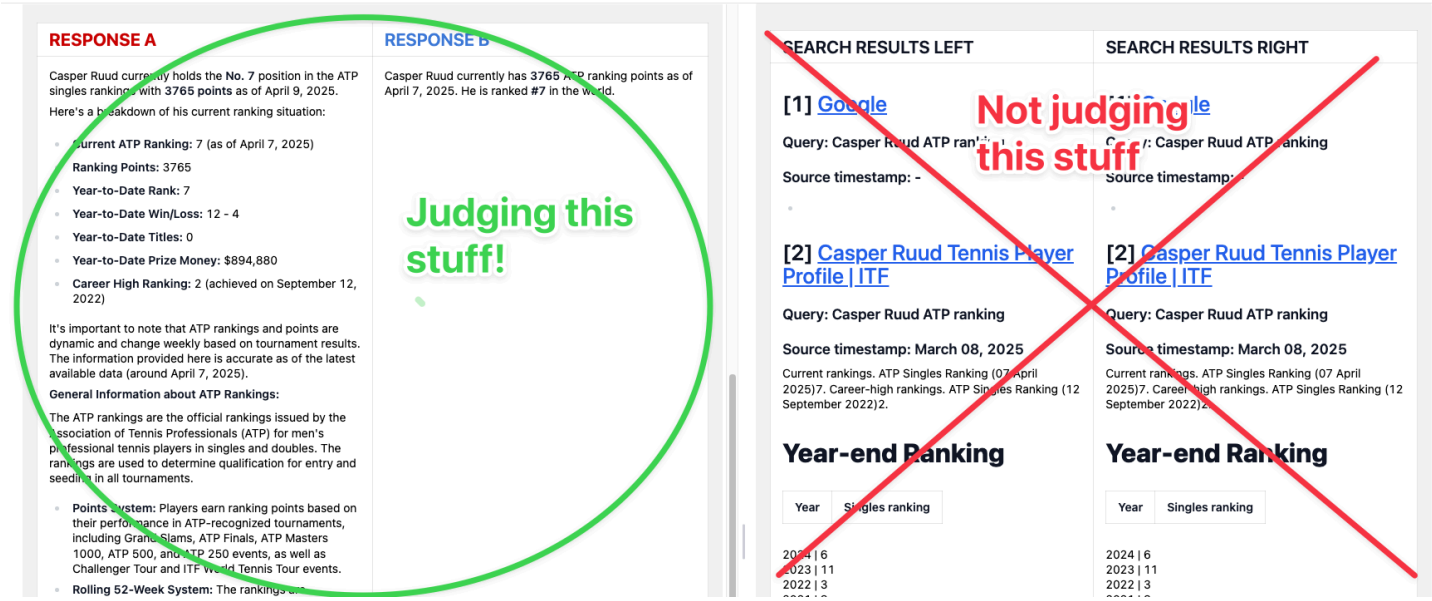


## **What is** Factuality?

Something that is a "Factual Claim" is any piece of a response that has to be grounded in real world information. Factual Claims are the information in the Response that users would expect to be correct and could either attempt to verify with additional research OR would expect to be true within the context of the conversation.

Among Claims, some may be more subjective than others. Objective Claims can be verified or refuted with specific information, while subjective Claims (e.g., opinions, recommendations) should be considered accurate unless they are unreasonable or there is reliable evidence that suggests many or most would disagree with the claim.

Objective Claims include but are not limited to:

- Numbers: dates, sports game results, rankings, distances, sizes, prizes, transportation and venue schedules, etc.
- Historical information, news, or any information about real-world events in the past, present and future.
- Information about people, locations, pieces of art, music, and other real-world tangible and intangible entities.
- Commercial products' information and specs.
- Information about the conversation itself - for instance, the model saying "this is a summary of…", "I have reduced this by 200 words….", "I have extracted all words starting with A…", etc.

# Ratings

## Overall Questions

1. Can the task be successfully rated? (Optional Selection)
   a. **Task cannot be rated**

   You will mark this if there is a situation with the prompt or response where you cannot rate the task. In general, we want to use "Prompt Cannot be Rated" (aka "won't do") for tasks where you **don't have the necessary information to assess** or tasks that **violate certain safety guidelines.**

   If you select this, you will have to mark the reason that makes it not ratable. Please see the FAQ Question for the "When should I mark a task as "Prompt Cannot be Rated"" reasons

2. Is the Prompt understandable?
   a. Mark **YES** if you can understand the prompt (even if you do not know how to answer it). If the Prompt is not in English (even if you do understand the language) or is incoherent, you will mark **NO**. We are *mostly interested in the Responses* in this project, so you will still rate the Responses even if you do not understand the Prompt. You will **not** have to rate which one is better, however.

## Response Questions

1. Is the response a **canned response**?
   **A canned response** is a predefined response that attempts to avoid the question, not answer it. Mark **YES** if the response you are evaluating falls under this category. If you mark **YES** , you will not need to mark anything else for this Response. If the model follows up with separate, factual information, however, please pick **NO**

   - **YES:** "I'm a language model and don't have the capacity to help with that."
   - **YES:** "I'm sorry, I don't understand what you're asking me to do. Can you please rephrase your question?"
   - **YES:** "That feature is not available right now."
   - **NO:** "I cannot answer prompts about how to use nitroglycerine as a bomb. However, I can tell you <facts about how it is used as a heart medicine"
   - **NO:** "As a large language model, I cannot fly to the moon. However, here is <some interesting history about moon travel>"

2. Does the response **attempt to convey any factual information**?

   **YES** if the response is likely to include information that must be grounded in the real world or in the conversation itself that must be accurate and verifiable. You can interpret claims in the context of the Prompt - ie, an answer of "No" to a yes/no question can be understood to be answering the question. **Still mark yes even if the response is a bad answer to the prompt but does contain factual information.**

   **NO** if the response is creative or otherwise does not contain any real-world information. If a response conveys no factual information, you will not need to perform the accuracy assessment.

3. Does this Response **hedge?**

   **YES** if the response provides some relevant information, or help on where to find it, but does not directly fulfill a user's request. Note, a Hedge response can be either accurate or not accurate.

   **NO** if the response directly answers the user's prompt (whether or not it includes extra information). Note - a "direct answer" can either answer positively ("who was the first US president" → "Washington") or can directly answer negatively ("who was the 75th US president" → "there have not been 75 US presidents; there have been …").

   **UNCERTAIN** if you cannot really tell whether or not the response is hedging

   Please take a look at the examples given here

4. Research the Factual Claims in the Response and then indicate the level of accuracy of the Response (**If the Response does contain factual information).** You must assess the accuracy of all Factual claims in the Response (unless you find one that is Inaccurate).

   Mark the first rating in the list that applies. As such, if you find one Inaccurate claim, you can mark the whole Response as Inaccurate and do not need to continue your research (this only applies if you find a fully Inaccurate claim; otherwise you must continue assessing, because you might find a "worse" claim).

Search in this order!

- Look up **key facts** first (ie, the points that are most important to answer the response)
- Then look up places where models may likely hallucinate (numbers, references, authors, links etc)
- After that, go through the remaining factual claims in the response.

You must still research the claims even if the response is a bad match to the prompt.

- **Inaccurate**: at least one Factual Claim is contradicted by reputable evidence (and no reputable supporting evidence) or is unreasonable
  - "George Washington was the third president of the United States"
  - "Unicycles are easy for babies to ride"

- **Unsupported**: at least one Factual Claim has no reputable evidence to consider (supporting or contradictory)
  - "Grigory Moskovich is a Russian who works in a bookshop"
    - This is almost impossible to fully disprove. There is no full database of every Russian who works in a bookshop. If you cannot find any evidence explicitly supporting this, you won't be able to explicitly disprove it, so it would be Unsupported.

- **Disputed**: at least one Factual Claim has a mix of supportive and contradicting evidence and this discrepancy cannot be reconciled
  - "Red wine is good for you"
    - There are reputable links both supporting and contradicting this; it is Disputed.
  - "Socrates was born in 469 BCE"
    - Some evidence says 470; others are vaguer in general. But there is *expert disagreement*
  - This is **NOT** for things where you "just aren't sure" or where you are unclear on what the claim really is. You should only pick Disputed when you can
    - 1. Clearly articulate the claim
    - 2. Find reputable evidence that both agrees AND disagrees with the claim as stated.

- **Accurate**: all Factual Claims are supported by reputable evidence (and no reputable contradictory evidence exists) OR are reasonable / widely agreed upon opinion statements.
  - "George Washington was the first President of the United States."
  - "Harvard is well regarded academically"

- **Can't confidently assess**: unable to understand, identify, or properly assess the claims in this response
  - Responses that are simply too muddled, poorly constructed, or non-grammatical such that you cannot assess their factual nature.
  - If a technical limitation blocks you from being able to assess a given claim (e.g., if geographic location blocks you from accessing a key URL, reports that cost thousands of dollars to view, etc)
  - Responses that are dependent on external knowledge that is inaccessible or private - internal business information, information about the bots themselves, or similar.
  - The Response is entirely irrelevant to the Prompt **AND** can't be determined to be true or false without that missing information.
    - Ex: Prompt: Why do bananas grow upside down?
    - Response: His first name is Thomas.

## 🚨 NEW SEVERITY INSTRUCTIONS - MUST READ 🚨

5. Identify the **most severe non-Accurate claim** in the response and rate its factuality level (if the response is completely Accurate, you'll skip this):
   - What you determine to be "the most severe non-accurate claim" may not always align with the overall accuracy hierarchy above -- severity completely depends upon the prompt/the user's needs. In other words, you're assessing the non-accurate claim that affects the answer the most.
   - In the event that two claims with different accuracy levels are equally severe (e.g. an inaccurate claim and an unsupported claim that both affect the quality of the response in the same way), you should then defer to the accuracy hierarchy to determine which to evaluate (Inaccurate > Unsupported > Disputed > Can't confidently assess).

Once you've identified the **most severe non-Accurate claim,** you'll rate that claim along a series of severity axes (steps 6 - 11).

6. 🎯 Centrality: **Is the most severe non-Accurate claim central to the prompt?**
   - That is, does the non-Accurate claim directly address the prompt and/or is its inclusion necessary to fulfill the user's intent?

     **Yes** - The most severe non-Accurate claim is central to the prompt.

**No** - The most severe non-Accurate claim is not central to the prompt.

7. 🤪 Obviousness (makes the model look silly): **Is the non-Accurate claim something that's easy to identify as non-accurate, even without researching it?**
   - Even if you don't know what the correct answer is yet, it's still obvious that the non-Accurate claim is wrong.
     - This would be obvious to *most* people -- put yourself in the position of, say, an intelligent 12-year-old.
     - If someone posted it on Reddit, would everyone instantly know it was wrong?
   - If it's "obvious," the non-Accurate claim is not an issue that would require higher or specialized education to understand (which, while insidious, might be more understandable to get wrong).
   - ⚠️ Only "Inaccurate" claims can be considered obvious. If the claim is "Unsupported," "Disputed," or "Can't confidently assess," it's not obviously wrong!

   **Yes** - The most severe non-Accurate claim is obvious.

   **No** - The most severe non-Accurate claim is not obvious.

8. ↔️ Distance from the truth for measurable statements (if the Response contains non-Accurate claims): **Is the non-Accurate claim far from the truth?**
   - Measurable means literally that – would you be able to say something like "what percent wrong it is?"
     - This is about things like monetary values, heights, years, dates, ages, etc.
   - Smaller errors, especially in numeric quantities, are more understandable and less likely to be detrimental than gross errors.
   - ⚠️ If the non-Accurate claim isn't measurable, or if it is "Unsupported," "Disputed," or "Can't confidently assess," select 'N/A' here.

   **Yes** - The most severe non-Accurate claim is measurable and far from the truth.

   **No** - The most severe non-Accurate claim is measurable, but it's not far from the truth.

   **N/A** - The most severe non-Accurate claim is not measurable.

9. 🌍 Sensitivity: Conspiracy Theories, Religion, Persons, Nations, etc. (if the Response contains non-Accurate claims): **Does the non-Accurate claim concern sensitive topics with direct real world impact?**
   - Claims concerning sensitive topics like conspiracy theories, religion, persons, nations, etc. with real world implications are critical to get right.

   **Yes** - The most severe non-Accurate claim concerns sensitive topics.

   **No** - The most severe non-Accurate claim does not concern sensitive topics.

10. 💸 🏥 YMYL (Your Money / Your Life): **Does the non-Accurate claim concern the health or money of the user?**
    - Claims concerning health/money topics with possible impact on the user are critical to get right.

    **Yes** - The most severe non-Accurate claim concerns the user's health or money.

    **No** - The most severe non-Accurate claim does not concern the user's health or money.

11. Rate the **Overall Severity** of the **most severe non-Accurate claim** present in the response on a scale of 1-5:
    - We're not putting hard rules on what you pick; we want to see what you think. But generally, something that is a major error on a central point regarding a sensitive topic would be more severe than a minor error on a non-central point.

    1 = the issue is least severe

    5 = the issue is most severe

12. Briefly explain your ratings including how you arrived at your overall factuality rating, how you determined the most severe non-Accurate claim, your Overall Severity rating, and whether there were other specific claims that

were more difficult to assess for that Response. If you answered "Yes" for any of the individual Severity axes, you should also explain those selections.

Your explanation should ideally note to what extent the model used the search results within the Facts Referenced section to form its response.

## Overall Comparison (Only for Understandable prompts)

If the prompt is understandable, you will assess which of the two Responses you think is better for the Prompt. You will be picking the Response that you think is better in general, though you should of course consider the **Accuracy** and **Information Quality**. The rating will be on a 7 point scale from A is {Much Better, Better, Slightly Better} to About the Same to B is {Slightly Better, Better, Much Better}. The two answers being equally bad would be “About the Same”.

You will need to **explain why you made your choice.** If one response has a failure, copy and paste that section of the response and note how it failed, like “wrong data point” or “controversial claim”.

**NOTE -** A canned response is not necessarily better or worse than a Response with factual statements, depending on the reasons behind the canned response. A canned response can be better if:

- The prompt asks about a private person that doesn't have any publicly available information and a wikipedia page under the exact same name.
- The prompt requests to quote content that is subject to copyright (song lyrics, books,etc.)
- The prompt contains offensive, controversial, or potentially dangerous/harmful content.
- The prompt requests something that the model cannot do.
- The other response is simply really bad!

## *Examples of identifying Factual Claims*

| and Response | Claims |
|---|---|
| **t:** Tell me about the movie Shawshank<br>ption<br><br>**nse:**<br>can help with that. The Shawshank<br>ption is a 1994 American drama film about a<br>who is wrongfully sentenced to prison for the<br>s of his wife and her lover. The film is a classic<br>t some viewers find it to be too long. | • The Shawshank Redemption is a 1994 American drama film *(objective)*<br>• The Shawshank Redemption is about a banker who is wrongfully sentenced to prison for the murders of his wife and her lover *(objective)*<br>• The Shawshank Redemption is a classic film *(subjective)*<br>• Some viewers find The Shawshank Redemption to be too long *(objective)*<br>  ○ *While "it is too long" would be a subjective claim, the claim in this response is that "some viewers find it to be too long" is an objective claim because it could be directly verified with evidence whether there are at least some reviews that indicate they found the film too long.*<br><br>Note that "Sure, I can help with that." does not contain any Claims. |
| **t:** What is Paris famous for?<br><br>**nse:**<br>/ of Paris is the capital of France and located<br>e-de-France region. It is a beautiful city with a<br>tory and culture and also considered as the<br>mantic city in the world. Paris iconic | • Paris is the capital of France and located in the Île-de-France region (*objective)*<br>• Paris is a beautiful city with a rich history and culture and also considered as the most romantic city in the world *(subjective)* |

| | |
|---|---|
| rks, such as the Eiffel Tower and the Louvre n, are known around the world. | • The Eiffel Tower is in Paris *(objective)*<br>• The Louvre Museum is in Paris *(objective)*<br>• Paris' landmarks are known around the world *(subjective)* |

## *How to Research*

- Unless you are 100% sure that something is factually accurate, perform a web search to **find evidence that supports the factual claims** in the response.
- You are researching **every factual claim** (in each answer!) until and unless you find a fully Inaccurate claim.
- Search in this order!
  - Look up **key facts** first (ie, the points that are most important to answer the response)
  - Then look up places where models may likely hallucinate (numbers, references, authors, links etc)
  - After that, go through **all** remaining factual claims in the response.
- Use a search query that is either:
  - A query **similar** to the user prompt.
  - A query that looks for a **specific relevant fact** for the response.
  - A query that confirms a **specific relevant fact**.

## Reputable Sources

- Use Reputable Sources
  - Reputable sources are those from trustworthy websites or creators with a high degree of expertise, experience, and/or authority for their content. These sources generally provide information that makes it easy to learn about their background, credentials, or process for producing their content. When researching important or sensitive claims that have the potential to be harmful if inaccurate, you should have an even higher standard for considering a source to be reputable.
  - It is best to look at multiple sources for agreement when researching evidence. When sources disagree with each other, it is sometimes possible to reconcile them based on factors about one source that may make it more trustworthy than another source, such as:
    - Freshness: If a response talks about information that may change over time, sources that were created more recently can be considered more reliable to have more up to date information. If a webpage is extremely outdated relative to how time-sensitive a particular claim is, then it might not be a reputable source for verifying that claim even if the website as a whole is considered trustworthy and authoritative.
    - Authority / primacy: Sources that are the primary, definitive, or original source for certain information can be considered more reliable than secondary or alternative sources.
    - Prompted evidence: If a prompt asks for information from a particular source, evidence from that source should be considered authoritative.
    - Topic context / nuance: Sources may disagree based on different assumptions or context associated with the information. Sources most aligned with the Prompt and Claim in the response can be considered most appropriate.
  - ✅ **Examples of Reputable Sources**
    - Wikipedia
    - High-quality online publications like ProPublica
    - Academic institutions
    - Official organizations (e.g. CDC, WHO)
    - News sites
    - Books
  - 🛑 **Examples of sources to avoid**
    - NEW! The “AI Overview” at the top of a Google search result
    - Social media sites (e.g. Facebook, LinkedIn, and Twitter)
    - Blog posts
    - Sites with fun fact lists
    - Opinion sites/Q&A sites (e.g. Reddit, Quora)
  - ☠️ **You may not “Fact Check” with any LLMs such as ChatGPT or any other bot.**

- You can potentially use them as a resource to flag items in, eg, long lists, that might be possibly wrong, but you **must** confirm any information with an actual reputable source.

# FAQ

## What should my comments look like?

- Your comments should always be the requested length (often this is 2-3+ sentences). Beyond that, *every comment* you write should make clear that (a) **it is for that task** and (b) **that a real human wrote it, not a robot copying from a template.**
- This helps prevent errors, gives us a little insight into what you're thinking, and makes it clear you're not on autopilot. It can also help keep you more engaged!
- In essence, if your comment is something you could apply to many different tasks, it's too generic and you shouldn't submit that comment (this is pretty much universal on the platform 🙂 )
- Train of thought of *which aspects* made you rate *this task* in *this specific way* will go far! What details caught your eye on this task? Which attributes were interesting? Let your own opinions shine through in each comment.
- Provide sources for your fact-checking - **especially** for any claims you are saying are not accurate (but in general - provide sources).
  - While this will often be web links, if you are using the **search results** from a model to prove/disprove some part of an answer, say that you used them!
- Check here for some examples of good and bad comments!

## Are these models up to date?

- Most of these models have access to real-time information and can access URLs. For example, if you ask them information about something that happened last week, or if you ask them something that depends on a specific website, they can search the web to do so. (If a model **says** it has a cutoff date, however, you should treat that as a true statement).

## Do I still rate the Responses even if they are a bad match to the Prompt?

- Yes - still rate them! We are primarily interested in the truth of the factual claims, even if they are unrelated to the Prompt. If the mismatch is so severe that you cannot actually tell whether the Response is correct, you may mark it as "Can't Confidently Assess". For instance, if a Prompt was "Do you like unicorns?" , and the answer was "His first name was Thomas". That sentence makes a claim, but can't be judged true or false without a related prompt.

## Do I still rate the Responses even if the Prompt is nonsense or not in English?

- Yes, still rate them! We are primarily interested in the truth of the factual claims. Most factual claims can be judged without a prompt; for the cases when you cannot, you can mark "cannot accurately assess". You will mark that you cannot understand the Prompt and will **not** need to judge which response is better when there is not a Prompt that is understandable.

## (New Example Added) What do I do with time-dependent searches or time-specific answers?

- Time-dependent : questions whose answers change with time ("who won this year's super bowl?"
- Time-specific: "As of June 18, 2022…."
- [All of the below is for evals where you are shown a prompt and a response. If you are chatting in real time with a model, you know when you sent the query! That's when you are judging it as]
- If a prompt, context or response gives a specified time, judge the information using that time as context.
  i. Ex: Whole response says. "As of May 15, 2023, Chris Evans is not married. He has not announced any wedding plans."
  ii. This would be ACCURATE, because it was accurate as of May 15, 2023 (although he is now married).
  iii. Ex: Prompt says "When was the most recent XYZ game?" with a context date of August 10, 2024.
  iv. Response says "The most recent XYZ game was on ABC date". If that was correct on August 10, 2024 it is ACCURATE, even if it would not be true at the time you rated the task.
- If **no** time is specified in any of the prompt, prompt context, or response, use "today's" date (ie, the date you are doing the work or maybe within a few days just to account for how long it might take to get to you. This might make some answers

Inaccurate that might have been accurate when they were created. That's ok.

- If a time is given in one Response but not the other, do **not** assume that the other Response must have the same time context.
- In order, respect the given dates as to when a Response was generated:
  i. Any date the Response gives ("as of today, June 2, 2023…")
  ii. Any date given in the context for the Response (that **should** be when it was generated)
  iii. Any date given in the prompt
    1. Note - a prompt could be from much earlier than a Response was generated. If the prompt says "as of today, june 2, 2023" but the model answer is from January 2024, consider how the answer would be phrased and what would be correct in January 2024.
  iv. "Today"

## Do I need to fact-check links?

- Yes 🙂
- If a model gives you a link, you expect that:
  i. The link truthfully is what the model says it is (ie, the model doesn't say "this youtube link is for a Ted talk" and then the link goes to anime videos)
  ii. The link works (it resolves, doesn't 404, etc).
- As such, a bad, misdirected, or broken link is a truthfulness issue (central or not depends on how the link is presented).
- CAVEAT: Links underneath images are ONLY LINKS TO THE IMAGE SOURCE. They are not references for the information in the response. As long as they work, they are fine.
- CAVEAT: "links" in the form of **[† Source]** aren't real links; you can ignore them
- CAVEAT: if the link has quotation marks (") at the end and it doesn't work when you click it, please try the same URL **without** the quotation marks before rating it bad.
- CAVEAT: links in the form of "googleusercontent.com" will not work; you can also ignore those.
- NEW, CAVEAT: You can IGNORE the "invalid URL removed" for truthfulness (though you can rate it down for quality).
  i. The **model** generated / produced a URL that a separate system determined was false. So the link was removed (there may be other processing later but you are seeing this interim step). The need to judge the validity of the URL is removed from your hands (the other system already did it), it is only interesting in a general quality level that there is now a sort of blank hole where a URL was supposed to be. But the model itself wasn't "lying" when it said "here is a URL"; it is post processing that removed it.

## What do I do with searches for rapidly changing time-dependent searches (weather, prices, etc?)

- For searches for things that you would expect to change very rapidly (weather, prices, airline prices, stocks), start with the same general time/date information from the above question
- NEW - On some projects, you **<u>may</u>** have access to search results the model received at the time it generated its answer. You can consider a Response Accurate if its answer matches that search information, and the search information is from an accurate source.
  i. Example - if the prompt was for weather but was from a few days ago, and the search results contain weather information from weather.com for that date, you can trust that, and the Response can be marked Accurate if it matches that search snippet.
- Within that time range, however, you can mark responses as Accurate if they seem "roughly correct", because we would expect there to be multiple reasonable weather forecasts for a given day; airline prices can change every minute, and Amazon and Walmart may have different prices for the same item.
- If you absolutely cannot find information about the weather/price/time of something on a given day, you may mark it Cannot Assess (presuming something else in the response is not Inaccurate, Disputed, or Unsupported!)

## How do I interpret the "Search results" when they are present?

- The search results reflect additional information the model received when it was generating its answer beyond its normal training data.
- The model **only has access** to the *<u>literal text</u>* that we are showing you as its search results - they do **not** have access to any other information on the website. We provide links only as evidence of where the information came from and to help you judge whether the source is trustworthy.
- If the source is trustworthy, you can presume that the information in the snippet text is accurate.
- This is expected to be most useful for rapidly changing information, like airfares, stock prices, weather forecasts, etc.
- For responses that don't require immediate search results, you should generally stick to your standard research to confirm / deny factuality 🙂
- It's fine for the model's response to be associated with only a few search results and for the model to rely mainly on latent knowledge.
- If you do depend on a search result to prove/disprove something a model said, please state this in your comment!

- UPDATE :"citations" in the form of **[† Source]** are not references to the search results (they'll generally be from a model without search results at all listed)

## If the model is asked to summarize, am I fact-checking the information in the summary?

- No 🙂
- You are fact-checking the **claims** the model is making. In any response of the style "The article/story/book says that : <stuff>", the claim is that the article/story/book *says those things.* There is no implicit claim that the information itself is true.
- You **may** however be fact-checking that the summary is an accurate representation of what the model is being asked to summarize, **if the model says** "Here is a summary of that article: <summary>".
- UPDATE - this is similar for rewrites or closed QA. In those cases, you are not checking whether it is true in a global sense. You may be checking (more for Closed Q&A) whether it is true *based on what was in the source text.*
  i. Ex: "Based on this article, why is the earth flat?"
  ii. Ans: "Based on the article, the earth is flat because a turtle sat on it" → the truth of this depends on what the source article said. Clearly it wasn't a turtle that actually sat on the Earth to make it flat.

## Are links claims?

- Yes 🙂
- If a model gives you a link, you expect that:
  i. The link truthfully is what the model says it is (ie, the model doesn't say "this youtube link is for a Ted talk" and then the link goes to anime videos)
  ii. The link works (it resolves, doesn't 404, etc).
- As such, a bad, misdirected, or broken link is Inaccurate.
- CAVEAT: Links underneath images are ONLY LINKS TO THE IMAGE SOURCE. They are not references for the information in the response. As long as they work, they are fine.
- CAVEAT: "links" in the form of [† Source] aren't real links; you can ignore them (the model did create them, but we simply can't evaluate them interestingly)
- CAVEAT: if the link has quotation marks (") at the end and it doesn't work when you click it, please try the same URL **without** the quotation marks before rating it bad.
- CAVEAT: links in the form of "googleusercontent.com" will not work; you can also ignore those.
- CAVEAT: You can IGNORE the "invalid URL removed" for truthfulness (though you can rate it down for quality).
  i. The **model** generated / produced a URL that a separate system determined was false. So the link was removed (there may be other processing later but you are seeing this interim step). The need to judge the validity of the URL is removed from your hands (the other system already did it), it is only interesting in a general quality level that there is now a sort of blank hole where a URL was supposed to be. But the model itself wasn't "lying" when it said "here is a URL"; it is post processing that removed it.

## When should I mark a task as "Task Cannot be Rated" (rather than Skip)?

- In general, we want to use "Task Cannot be Rated" (aka "won't do") for tasks where you **don't have the necessary information to assess** or tasks that **violate certain safety guidelines.**
- **THIS LIST IS SLIGHTLY DIFFERENT THAN THE EVAL LIST :** Since we are primarily concerned with the factual information in the Responses, you will still largely do factuality rating even in cases where you cannot understand the Prompt. You will still mark cannot be rated when:
  - Contains PII: if a prompt or response contains PII (personally identifiable information; like a full name, or name + address for someone the person is writing a letter to; general names of celebrities or other common figures are fine).
    - If the information is enough to **identify a person**, then it is PII
  - Response Incoherent / gibberish: If you cannot understand the Response
  - Response Non-English: if all or most of one or both responses are not in English
    - This includes answers "in English" that are translations from another language, since you would need to speak that language to actually assess the claim that the Response is a fair translation.
  - Refers to media uploaded by user: If a prompt seems to have contained a link/file/image that you don't see, such that you cannot assess the response.
  - Depends on unknown user location: If a prompt depends on knowing the user's location and the user hasn't shared a location to use and there is no location given in the context.
  - Mid-conversation turn: prompts that are clearly in the middle of a conversation ("can you explain that further?")
  - …anything else where you **cannot** assess the prompt/response pair (in depth math knowledge, detailed legal/medical/coding knowledge on a general knowledge project; other unintended things that may crop up)

### Can you expand what you mean by “cannot assess the prompt/response pair”?

- While you will be expected to try to evaluate most Responses, if a Prompt is ~”out of scope” and is not something a general worker could assess even with web research, you can mark that as “Prompt Cannot be Rated” with category “Other” and submit. Leave a comment when you do this.
  - Don't use this options for information you are simply less familiar with, but only for cases where you can't "look things up" to verify an answer
  - Even if you happen to be the type of expert that could actually rate it, please still mark it cannot be rated!
- ❌ Don’t Rate This: "This code will do X"
  - If someone cannot code, there is no way to assess this statement without literally trying to run it (which is not web research)
- ❌ Don’t Rate This: "Here is a good chemical synthesis of maltodextrin: <multiple synthetic steps>"
  - Even advanced chemists would have difficulty validating a given chemical synthesis; advanced PhD level work is not web research.
- ✅ Try to Rate This: "The change in Gibbs free energy can be calculated with: $\Delta G = \Delta H - T\Delta S$"
  - You may not know anything about this, but you can fairly easily look this up
- ✅ Try to Rate This: "The Riemann curvature tensor (named after Bernhard Riemann) can be used to express the curvature of Riemannian manifolds."
  - You may not know anything about this, but this is the first line on Wikipedia about "Riemann curvature tensors" and is a simple yes/no fact you can at minimum try to look up.

### What should I do when the tables in the response are overlapping?

- The models are fine (don’t rate them down!); unfortunately, our interface has some issues with width when the two responses both have tables. It tries to give them both all the room they need, so they overlap. You can zoom out (make the text smaller) in your browser; this can make them stop overlapping. If you find that the text is then too small to read, you can cut and paste the table text into a word doc or textpad. If none of that works for you, please just skip the task. You don’t need to call these out (though of course if there are other new, strange, problematic formatting errors, please do let us know!)

### I don’t have a “Which response is better?” button, what should I do?

- You need to hit YES to “is this prompt understandable” to make that button show up. Please double check that. In a very small number of cases, if you did hit Yes, you will just need to refresh your page, but that should be <1% of cases.

### If all claims are Accurate, what should I write in the “Briefly explain your rating for Response A/B” box?

- If all claims are in fact accurate, your explanation for that side can be short, but it should make it clear that it is (a) **for that task** and (b) **that a real human wrote it, not a robot copying from a template. - [NOTE - these are lines for Responses that would be really short. If there are more facts, you would write a longer comment]**
  - ✅ `All ridiculous Great Pumpkin / Peanuts facts confirmed via wiki! https://en.wikipedia.org/wiki/Great_Pumpkin`
  - ✅ `It is clearly true that grass is green and the sky is blue`
  - ✅ `Everything about George Washington checks out...even the bit about his dog! https://www.whitehouse.gov/about-the-white-house/presidents/george-washington/`
  - 🔴 `All facts verified`
  - 🔴 `Everything confirmed`
  - 🔴 `All claims are accurate`
- This helps prevent errors, gives us a little insight into what you’re thinking, and makes it clear you’re not on autopilot.
- In essence, if your comment is something you could apply to many different tasks, it's too generic and you shouldn't submit that comment (this is pretty much universal on the platform 🙂 )
- Note - your OVERALL side by side comment must still conform to all length + content expectations

### Can I see some more examples of “hedging”?

- Sure - down here: Examples of hedging

### **What do hints look like?** What do I do if the hint is wrong?

- Hints look different depending on the prompt. For example, if the prompt is a math problem, then the hint will show the correct answer. For other prompts, the hint might explain the desired properties of the response (e.g. start with something, push back on something, highlight something, avoid mentioning something, refer to something).
- Generally, hints are accurate and can be considered as the ground truth. However, always use your own best judgment. If you strongly believe a hint is wrong, then go with your own judgment.

# Examples

## Examples of Accuracy ratings

| Prompt | Response | Rating Explanation |
|---|---|---|
| What type of bicycle is best to get my 5 year old who is a beginner and will be riding on trails and dirt often? | A unicycle would be fun and engaging for your 5 year old to ride trails with. | **Inaccurate**<br><br>The Response contains a subjective claim. While it may be hard to find evidence that directly disagrees with the Response, you could infer this is an unreasonable claim from evidence that unicycles are more challenging than other types of bicycles to ride, particularly on trails and dirt. |
| When did Leonardo da Vinci paint the Mona Lisa? | The Mona Lisa was **painted between 1503 and 1519** by Leonardo da Vinci. The painting is a portrait of Lisa del Giocondo, the wife of a Florentine merchant. The Mona Lisa is currently on display in the Louvre Museum in Paris, France. The Mona Lisa is considered to be one of the most famous paintings in the world and is often regarded as a masterpiece of Renaissance art. | **Disputed**<br><br>The time period in which the Mona Lisa was painted is still unknown and reputable sources suggest different dates. The claim that the Mona Lisa "was painted between 1503 and 1519" is disputed.<br><br>The last sentence is mostly a subjective claim that research can verify that it is very reasonable. But also note that it includes an objective claim as well that "the Mona Lisa is Renaissance art". |
| What are some reviews of French Louie restaurant on Atlantic Avenue? | "The food was amazing! I had the steak frites and it was cooked perfectly. The service was also excellent." - Kseniya K. | **Unsupported**<br><br>The Factual Claim is that a person ("Kseniya K.") wrote this quote about the restaurant "French Louis", however there is no evidence of this review existing when searching on the web. |
| Who won the last NBA championship? | The first humans to step on the moon were American astronauts Neil Armstrong and Buzz Aldrin as part of NASA's Apollo 11 mission on July 20, 1969. | **Accurate**<br><br>Although the Response has nothing to do with the Prompt and is irrelevant and unhelpful, the information and Factual Claims in the Response are accurate. |
| Write a story about unicorns | A unicorn was walking in the forest when it came across a beautiful lake. | **There are no claims in the Response to check for accuracy** |

| | | |
|---|---|---|
| | The unicorn looked at its reflection in the water and realized that it was the only unicorn in the world. The unicorn was so sad that it cried, and its tears fell into the lake. The unicorn's tears turned the lake into a magical healing pool, and anyone who swam in the pool was cured of all their illnesses. | This is a fictional story and contains no Factual Claims. |

## Examples of hedging

| Prompt | Response | Hedge? |
|---|---|---|
| List me open PhD positions in Europe | "Absolutely! Here's a breakdown of a few reliable sites to find open PhD positions in Europe... | **Hedge** - although the response is helpful, it does not directly answer the prompt with a list of open positions. |
| Difference in sunset times for San Francisco and New York | There are a few factors that cause the difference in sunset times between San Francisco and New York... | **Hedge** - although the response is correct and there may be no one right answer without additional context, the prompt is not directly answered. |
| Give me cocktail bars near my place | Absolutely! Since I don't know your exact location, I'll focus on a couple of great cocktail bars in Mountain View, California... | **Hedge** - although the response is correct and there may be no one right answer without additional context, the prompt is not directly answered. |
| Is Jon Doe going to win the 2024 elections? | Elections are a complex topic with fast-changing information. To make sure you have the latest and most accurate information, try Google Search. | **Hedge** - although the requested answer would be a prediction and would likely need some caveats to be useful, the current response does not provide one as the prompt requests. |
| What are the results of the general election? | Elections are a complex topic with fast-changing information. To make sure you have the latest and most accurate information, try Google Search. | **Hedge** - although the model may not be able to answer because it lacks real time information, it does not directly answer the prompt. |
| When did the universe begin? | No one is certain, but many experts believe it is 13.8 billion years ago. | **Not a Hedge** - the model mentions uncertainty, but provides what it says is a widely accepted estimate that directly answers the question. |
| Did England win the soccer match tonight? | No, England has not won tonight. | **Not a Hedge** - although England may not have played or didn't win the response attempts to answer the user's question. |
| How many GPUs are currently on the market? | There are no new GPUs on the market in the last month. | **Not a Hedge** - regardless of whether or not it is correct, the response directly answers the prompt. |
| Tell me about Tome Hanks | Did you mean Tom Hanks? | **No Claims (not Hedge) -** the response asks a clarification question that cannot be assessed for accuracy. |
| What is the best ice cream? | Haagen-Daaz is the best ice cream. | **Not a Hedge** - this answers the question directly (though violates persona guidelines!) |
| | There are many opinions about what makes ice cream the best. Ultimately, the decision must be your own. However, here are some top rated ice creams... | **Hedge -** it doesn't say what the best ice cream is. |
| | Here are some top-rated ice creams... | **Hedge -** it doesn't say what the best ice cream is. It does not need "hedging language" to be a Hedge, please note. |

| | | |
|---|---|---|
| Where can I pet a real live unicorn? | You cannot pet a unicorn; unicorns are not real. | **Not a Hedge** - this directly addresses the request and says it isn't possible (it doesn't say the model can't or won't help; it directly addresses what was asked). |
| | | |

## Examples of Good and Bad comments

**Prompt:** What is the name of Frank Maxwell Andrews's child?

| Response A: | Response B: |
|---|---|
| Frank Maxwell Andrews actually had three children with his wife, Jeannette Andrews. Their names are:<br>Josephine (1914-1977)<br>Allen (1917-2008)<br>Jean (b. 1923) | Frank Maxwell Andrews and his wife, Jean Isabelle Hixon, had three children:<br>Josephine Andrews (1914-1977)<br>Allen Andrews (1917-2008)<br>Jean Andrews (b. 1923) |

**Good Comment:** According to Frank Maxwell Andrews' Wikipedia page (https://en.wikipedia.org/wiki/Frank_Maxwell_Andrews), both responses provide the correct names and birth/death dates for his children. The names of the children across the two responses are actually identical other than the fact that Response B tags the family name onto the children's names.

However, the same Wikipedia article states that he married Jeannette "Johnny" Allen, which is a different name from the one provided in Response B. I was able to locate the Wikipedia page for a "Jean Hixson" who was part of the United States Air Force. It's possible that these two individuals are being confused with each other in the response.

Response A is better due to this inaccuracy in Response B. It also corrects the prompt's implication that the man only had a single child, which is a slight bonus. It's only "better" rather than "much better" however since the other aspects of Response B are accurate.

– The comment not only highlights the factual claims made in the two responses and states which are Accurate/Inaccurate, but also explains the reasoning and the relevant evidence that was used when making these determinations. Notice how it clearly conveys the factors that worked in each response's favor and how the worker was weighing them against each other in their mind ("Response A is "better" because X, but because Response B still did well at Y Response A wasn't rated as "much better")

**Bad Comment:** Response A was accurate when providing the names while Response B wasn't, so Response A is better.

– This comment is vague and doesn't provide any clues as to which specific factual claim was accurate in Response A but not Response B. Prefer to be precise when discussing the accuracy of claims, especially when there's a meaningful difference relating to that claim between the two responses that sways the overall comparative rating.

**Prompt:** What was the opening price (USD) of VOO on April 2, 2024 as well as the highest and lowest prices? Keep your response brief and just state the figures, I don't need analysis.

| Response A: | Response B: |
|---|---|
| Information on VOO for April 2nd, 2024:<br><br>Opening Price: 476.33<br>Highest Price: 476.99<br>Lowest Price: 474.71<br>Based on available data, the opening price was close to the day's high, indicating a strong start. The low point was relatively close to the opening, suggesting limited intraday volatility. | Based on the available data:<br><br>April 2nd, 2024:<br><br>Opening Price: 469.33<br>Highest Price: 476.99<br>Lowest Price: 468.71 |

**Good Comment:** All three of Response A's figures match those present in the search results that were obtained from nasdaq.com. Response B also had access to the same figures from the same source in its accompanying search results but doesn't accurately quote them. It provides the correct Highest Price, but inaccurate values for the other two fields. Although Response A is

completely accurate, it does go against the instructions somewhat since it provides a little bit of analysis (the prompt indicates the user wants a more straightforward answer and doesn't need it). Providing the correct prices meets the most crucial request of the prompt though so I still rated it as better than Response B.

– The comment explicitly mentions that the search results were used when proving/disproving the claims across the two responses. It mentions each meaningfully relevant difference between the two responses, including how Response A conflicts with the instructions slightly by providing an analysis of the different price values. It also provides a justification for their overall comparative rating and why one response's issues were more significant than the other's.

Bad Comment: Response B's Highest Price figure is accurate, but its other two values are wrong. Response A is better, although it's too long.

– Although the comment conveys the main factors that justify the "Response A is better" rating, it has two issues:

- It discusses the accuracy of Response B's claims but doesn't mention the source that was used for fact checking. How exactly did you figure out that the Highest Price figure was accurate but the other values were wrong?
- The comment doesn't clarify whether the criticism that Response A was too long was inherently related to verbosity or whether it had more to do with instruction following. Don't leave us guessing; clue us into your train of thought!

Prompt: What is the name of utah's basketball team?

| Response A: | Response B: |
| --- | --- |
| Utah actually has two major basketball teams:<br><br>Utah Jazz: This is the professional team, playing in the National Basketball Association (NBA).<br>Utah Utes: This is the college team, playing in the NCAA Division I and the Pac-12 Conference. They are also known as the "Runnin' Utes."<br>If you were interested in a specific team, please let me know and I can provide you with more information about them. | The Utah Utes men's basketball team, also known as the Runnin' Utes. |

Good Comment: Both responses are right to mention the Utah Utes as one of Utah's basketball teams. I confirmed all of these claims across the two responses using the team's wikipedia page (https://en.wikipedia.org/wiki/Utah_Utes_men's_basketball). They are indeed also known as the Runnin' Utes and play in the NCAA Division I as well as the Pac-12 Conference.

However, Response A recognizes that the prompt could also be referring to a different team and clarifies this. Response B doesn't seem complete without this clarification as the other team is also a very important utah basketball team. I confirmed Response A's claim that the Utah Jazz plays in the NBA using the team's Wikipedia page. For this reason, Response A is better.

– This comment paints a complete picture regarding the accuracy of the claims made across the two responses, as well as how they were fact-checked (the sources that were used). It also discusses the direct rationale for why Response A was rated as being much better.

Bad Comment: Both responses are completely accurate, but Response A is more detailed so it's superior.

– Although this comment is entirely correct, it's incredibly generic and provides no evidence that the worker wrote it specifically for the task at hand. Both responses might be accurate, but how and with which sources was this confirmed? Why exactly was Response A's greater depth of detail relevant in this case? There's no explanation for why the extra detail was needed when answering this particular prompt.